

Shantanu Shastri

shantanushastri70@gmail.com | (412) 608-7945 | linkedin.com/in/shantanushastri

EDUCATION

Carnegie Mellon University	December 2026
Master of Information Systems Management (Focus on AI & Machine Learning)	GPA: 3.85
Manipal Institute of Technology	August 2022
Bachelor of Electronics and Communications Engineering	GPA: 3.99
<i>Relevant Coursework:</i> Generative AI, Distributed Systems, Data Structures, Algorithms, System Design	

SKILLS

Languages: Python, TypeScript, SQL, Java, C++, JavaScript, Bash/Shell
AI & Frameworks: Modern LLMs, LangChain, LlamaIndex, AI-Centric Systems, AI Agents, Evaluation Pipelines
Frontend & Data: React, Next.js, User Interfaces, FastAPI, Node.js, Relational Databases, PostgreSQL
Cloud & Deployments: AWS Deployments, Docker, Kubernetes, CI/CD, Apache Kafka, System Design

EXPERIENCE

BlackRock	January 2022 – August 2026
<i>AI Engineer Internship</i>	June 2026 – August 2026
AI-Centric Systems: Architected AI-centric systems end-to-end, integrating modern LLMs into real-world applications by engineering a robust multi-agent platform (RockAI ADK) to automate compliance workflows. Evaluation Pipelines: Built rigorous evaluation pipelines and defined metrics to measure AI agents and workflows, tracking data drift and latency securely to ensure continuous improvement and overall robustness. Modern AI Frameworks: Developed real-world AI applications using LangChain and LlamaIndex to process 3.5M records with a hybrid vector-graph backend, delivering high-quality features quickly.	
<i>Software Developer 2 - Associate</i>	January 2025 – July 2025
React & Next.js Integrations: Designed user interfaces utilizing TypeScript and React/Next.js, effectively deploying AI models into real-world apps to successfully translate natural language directly into JQL. End-to-End Data Workflows: Comfortably owned end-to-end data workflows by architecting scalable Spring Boot data frameworks that flawlessly automated 1.5M+ client touchpoints with high reliability.	
<i>Software Developer 1 - Analyst</i>	July 2022 – December 2024
SQL & Relational Databases: Demonstrated deep proficiency with SQL and relational databases by refactoring legacy PostgreSQL schemas on tables exceeding 50M+ rows, reducing query latency by 37%. AWS Deployments & CI/CD: Managed large-scale AWS deployments by building comprehensive CI/CD pipelines with integrated JUnit testing suites for highly secure, production-ready code delivery. Ambiguous Challenges: Tackled ambiguous challenges head-on with a bias toward action, engineering a sharded server environment with Zookeeper to boost system throughput to 10K+ RPS.	
Shunya	June 2020 – September 2020
<i>Data Science Intern</i>	
Real-World Impact: Kept up to date with new models by developing CV workflows on Linux, delivering immense real-world impact by fully automating and reducing manual grading time by 90%. Data Pipelines: Owned end-to-end data workflows to design data pipelines that dynamically recommended targeted practice, directly impacting end-user learning outcomes through data-driven decisions.	

PROJECTS

High-Performance Modular Booking Architecture
Cloud & Deployments: Designed and built end-to-end data pipelines and scalable backend systems using Apache Kafka, effectively handling bursts of 100,000+ real-time requests with robust AWS deployments.
Zulip Open-Source Engine Extensions
AI-Centric Systems: Architected AI-centric systems by building a high-throughput message engine in TypeScript/Node.js to seamlessly manage stateful data streams and LLM-driven chat summaries.