

Shantanu Shastri

shantanushastri70@gmail.com | (412) 608-7945 | linkedin.com/in/shantanushastri

EDUCATION

Carnegie Mellon University	December 2026
Master of Information Systems Management	GPA: 3.85
Manipal Institute of Technology	August 2022
Bachelor of Electronics and Communications Engineering	GPA: 3.99
<i>Relevant Coursework:</i> Machine Learning, Deep Learning, MLOps, System Design, Data Structures	

SKILLS

Languages: Python, Java, GoLang, Scala, SQL, C++, TypeScript, Bash/Shell
ML & AI: Supervised/Unsupervised Learning, End-to-End ML Pipelines, Feature Engineering, Model Training
Frameworks & Processing: PyTorch, TensorFlow, Scikit-learn, Apache Airflow, Pandas, Docker, Kubernetes
Deployment & Monitoring: MLflow, Prometheus, Grafana, Milvus, Neo4j, FastAPI, CI/CD, Model Evaluation

EXPERIENCE

BlackRock	January 2022 – August 2026
<i>AI Engineer Internship</i>	June 2026 – August 2026
End-to-End ML Pipelines: Implemented robust ML pipelines for data ingestion and feature engineering, designing a learning engine with GPT-4 classification (92% precision) to support financial decision-making. Low-Latency Inference: Integrated ML models into production systems, building a hybrid vector-graph pipeline (Milvus, Neo4j) to rank remediation paths with <200ms low-latency inference for 20K exceptions/week. Monitoring & Iteration: Monitored model performance and data quality in production via closed-loop agentic capture, actively debugging issues and improving recommendation accuracy by 15% over 8 weeks.	
<i>Software Developer 2 - Associate</i>	January 2025 – July 2025
Model Training & Fine-Tuning: Designed and deployed generative ML models translating natural language to JQL, fine-tuning Llama-3 (LoRA) on a 40K+ corpus to achieve 94% Exact Match in evaluations. AI-Powered Services: Maintained AI-powered services by deploying a Retrieval-Augmented Generation (RAG) pipeline into production, ensuring reliable operation and cutting query formulation time by 91%.	
<i>Software Developer 1 - Analyst</i>	July 2022 – December 2024
Supervised Learning: Applied supervised learning techniques to build an XGBoost model, experimenting with and optimizing features to achieve 88% precision, deployed reliably via FastAPI and Docker. Data Processing Frameworks: Engineered scalable data processing frameworks via Apache Airflow to ingest and clean multi-terabyte transactional data securely, adhering strictly to financial compliance requirements. ML Software Engineering: Applied software engineering best practices to codebases by building Azure CI/CD pipelines with comprehensive integration tests and rigorous code reviews for full ML components.	
Shunya	June 2020 – September 2020
<i>Data Science Intern</i>	
Unsupervised Learning: Leveraged unsupervised learning techniques to train K-means clustering models on multidimensional assessment data, dynamically engineering features to recommend targeted practice. Computer Vision Deployment: Developed and deployed GPU-accelerated PyTorch CV models on Linux to analyze answer sheets, reliably automating evaluations and reducing manual grading time by 90%.	

PROJECTS

FinTech Recommendation & MLOps System
Full Lifecycle ML: Developed a hybrid FinTech recommender (LightFM, SVD++) with a 4-hour retraining loop, evaluating data quality, managing version control, and monitoring operational performance via MLflow.
Real-Time Generative AI Product Agent
Performance Optimization: Deployed real-time AI services using Meta's FAISS for high-speed vector retrieval, optimizing the end-to-end ML pipeline latency to ensure rapid, low-latency inference in production.